

Agents

- Rational:** Maximally achieving goals (actions that maximize utility function)
- Reflex Based:** Chooses action based on current percept (no future consideration)
- Goal Based:** Chooses action based on consequences (model of how the world reacts)
- Utility Based:** Goal based with trading off of multiple goals and uses probabilities

Search

- Def:** Possible states, Successor function $f(n) \rightarrow (n', \text{action}, \text{cost})$, start and goal state
- Complete:** Guaranteed to find a solution if one exists
- Optimal:** Guaranteed to find the least cost path
- Properties:** n= number of states, b= maximum branching factor, C^* = optimal cost, d= depth of shallowest solution, m= max depth, ϵ = min cost of all actions
- Conformant Planning:** Set of actions that always work (sterilizing surgical gear)

Blind Search

- DFS:** Fringe uses a Stack, complete iff finite, not optimal, time: $O(b^m)$, space: $O(bm)$
- BFS:** Fringe uses a Queue, complete, optimal (constant), time and space: $O(b^d)$
- IDDFS:** Fringe uses a Stack, complete, optimal (constant), time: $O(b^d)$, space: $O(bm)$

Heuristic Search

- Heuristic** $h(n)$ = An estimate of how close a state is to a goal
- Admissible:** Always an underestimate to the true lowest cost
- Consistent:** Always $h(n) \leq h(n') + \text{stepCost}(n')$ where n' is a neighbor of n
- Best First:** Fringe uses a PriorityQueue with cost fuction for each node
- Uniform Cost:** Best First with $f(n)$ = sum of edge costs from start to n (explores increasing contours), complete, optimal, time and space: $O(b^{\frac{C^*}{\epsilon}})$
- Greedy:** Best First with $f(n) = h(n)$ (suboptimal goal is common)
- A*:** Best First with $f(n) = g(n) + h(n)$ with $g(n)$ = sum of costs from start to n
- IDA*:** Depth bound is now $F_{limit} = h(start)$, prune if $f(n) > F_{limit}$, $F'_{limit} = \min(\text{pruned nodes})$, uses space of DFS, time depends on # of unique F values
- Beam:** Best First with $|Fringe| = K$, not complete, time: $O(b^d)$, space: $O(b + K)$
- Hill Climbing:** Always choose best child (Beam Search with K = 1)
- Tabu:** Keep fixed length queue of states to not visit again (use with hill climbing)

Stochastic Search

- Hill Climbing++ Restarts:** Generate random state when plateaued
- Hill Climbing++ Walk:** With prob p move to the neighbor with largest value, with $(1 - p)$ move to a random neighbor
- Hill Climbing++ (Both):** Greedy move, random walk, or random restart
- Simulated Annealing:** Pick a random neighbor and calculate the change in ‘energy’ or objective function δ , if it is positive then move to that state. Otherwise, move to this state with probability $e^{\frac{\delta}{T}}$ where T is decreased as the algorithm runs longer. High T \rightarrow probability of bad move is higher and vice versa
- Genetic:** Start with a population of random states, use an evaluation (fitness) function, produce next generation using random selection / crossover / random mutation
- Gradient Descent:** Move in the direction of the gradient at each step

Constraint Satisfaction Problems

- Def:** Goal test is a set of constraints over the state’s variables $x_i \in D_i$ or D
- Constraint Graphs:** Nodes are variables, (multi)edges show constraints
- As Search Problem:** States defined by the values assigned so far, initially empty, Successor function assigns a value to an unassigned variable, and the Goal test checks to see if the current assignment is complete and satisfactory
- Improvements:** Fix ordering with variable assignments, check constraints as you go
- Forward Checking:** Cross off values that violate a constraint when added to the existing assignment (Immediate neighbors and fail if the set of possible values is empty)
- Constraint Propagation:** If X loses a value, neighbors of X need to be rechecked